

The Wumpus World: A Demonstration of Probabilistic Decision Making

Bryan Bugyi
Advised by Jonathan Weisbrod

Rules of Wumpus World

Description

- A Wumpus World is a cave consisting of rooms connected by passageways.
- Lurking somewhere in the cave is the terrible wumpus, a beast that eats anyone who enters the room.
- Some rooms contain bottomless pits that will trap anyone who wanders into these rooms.
- The only mitigating feature of this bleak environment is the possibility of finding a heap of gold. The agent's goal is to find the gold without falling into a pit or being eaten by the wumpus.

Environment

- The environment is a 4×4 grid of rooms.
- The agent always starts in a corner square.
- The locations of the gold and the wumpus are chosen randomly, with uniform distribution, from the remaining squares.
- In addition, each remaining square has a 0.2 probability of being a pit.

Sensors

- If a square is adjacent (not diagonal) to a pit, the agent will sense a "Breeze" when at this square; if the square is adjacent to the wumpus, the agent will sense a "Stench".
- The agent is only aware of the contents of the square that it currently resides in.**




Figure 1 : An example wumpus world environment

Important Definitions

$$E = \{\mathbf{x} \mid \mathbf{x} \text{ is a square in the environment}\}$$

$$V = \{\mathbf{x} \mid \text{The agent has visited } \mathbf{x}, \mathbf{x} \in E\}$$

$$N(A) = \{\mathbf{x} \mid \exists \mathbf{a} \in A \text{ where } \mathbf{x} \text{ is adjacent to } \mathbf{a} \text{ or } \mathbf{x} = \mathbf{a}\}$$

Wumpus Probability

Let $T = \{\mathbf{t}_1, \dots, \mathbf{t}_m\}$ represent the set of all squares that the agent has visited where a "Stench" sense was *not* detected.

Let $S = \{\mathbf{s}_1, \dots, \mathbf{s}_n\}$ represent the set of all squares that the agent has visited where a "Stench" sense *was* detected.

We can now use S and T to derive the set of all squares where the wumpus could possibly be found; we shall call this set W .

$$W = \begin{cases} N(\mathbf{s}_1) \cap \dots \cap N(\mathbf{s}_n) - N(T) & \text{if } S \text{ is non-empty} \\ E - N(T) & \text{otherwise} \end{cases}$$

Final Result

Therefore, by the axioms of probability, we can deduce that

$$P(w_{\mathbf{x}}) = \begin{cases} \frac{1}{|W|} & \text{if } \mathbf{x} \in W \\ 0 & \text{if } \mathbf{x} \notin W \end{cases}$$

where $P(w_{\mathbf{x}})$ represents the probability of a wumpus being located at square \mathbf{x} .

Utility Function

When deciding on which square to move to next, the agent will consider the square \mathbf{x} only if $\mathbf{x} \in V'$ where $V' = N(V) - V$.

The agent will consider two factors when deciding on the best next square: *death probability* and *distance*, where the former takes precedence over the latter. If all squares considered have a death probability of 0.5 or higher, the agent will forfeit.

Once decided on the next square to travel to, the agent must find the shortest path to said square...

Graph Structure / Shortest Path

After the next square is chosen, the agent models a portion of the environment using a graph-like structure where the nodes are $\{\mathbf{x} \mid \mathbf{x} \text{ is the target square or } \mathbf{x} \in V'\}$.

Figure 2 shows an example of how this structure can be visualized using an undirected graph.




Figure 2 : Visualization of the graph structure used to determine the shortest path from (2, 3)—the agent's location in figure 3a—to (3, 4)—the agent's location in figure 3b.

The agent then calculates the shortest path (shown in red in figure 2) with the help of Dijkstra's algorithm.

Running Examples

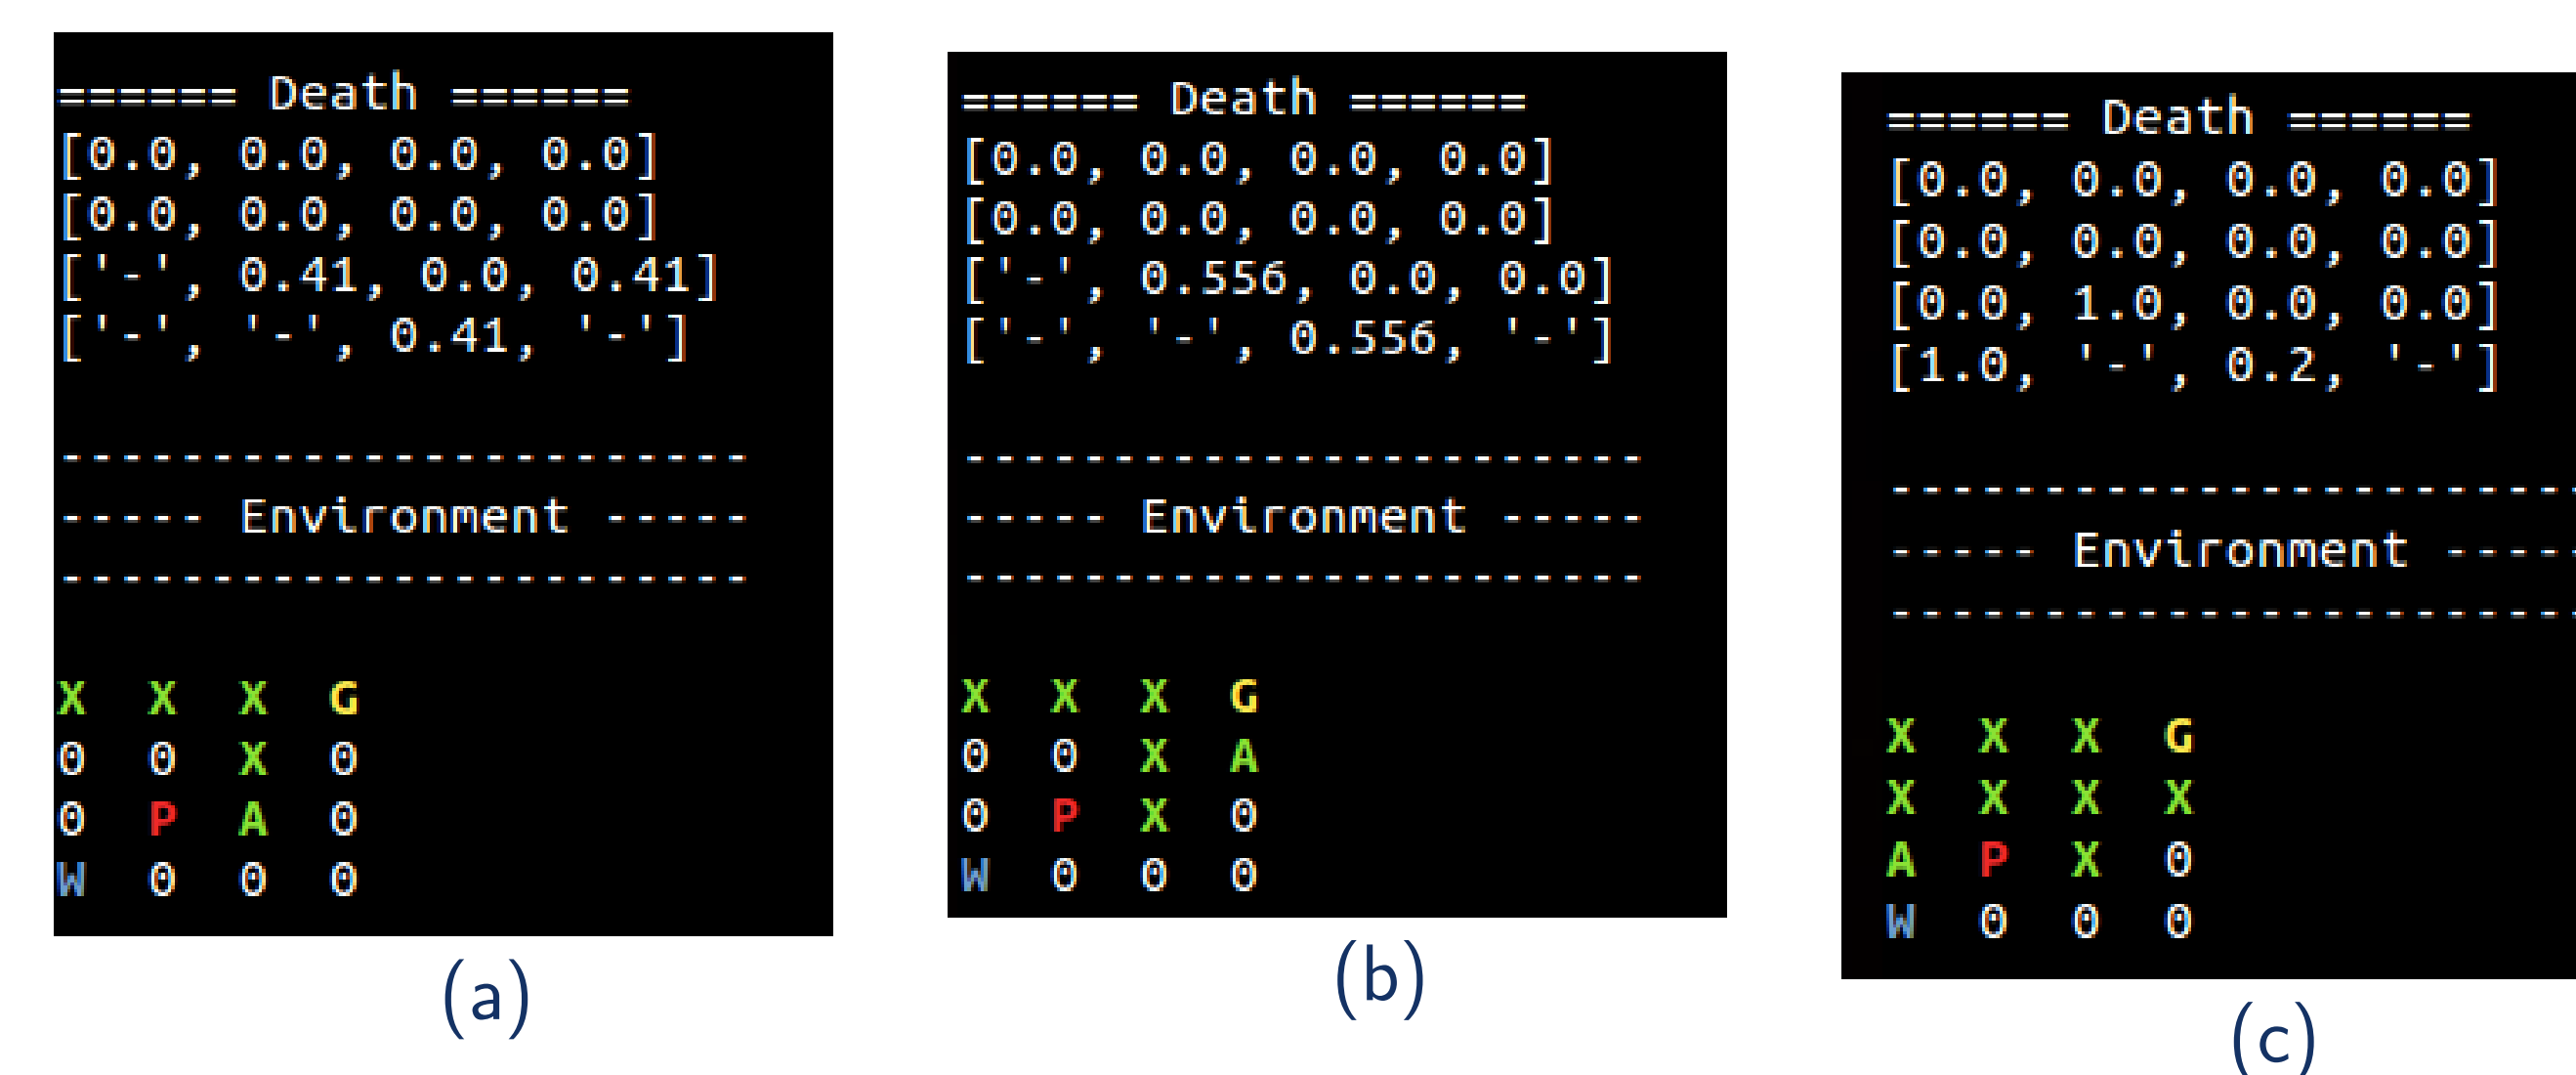


Figure 3 : A test run showing WW environment along with the agent's perceived death probabilities $[P(w_{\mathbf{x}}) + P(p_{\mathbf{x}})]$

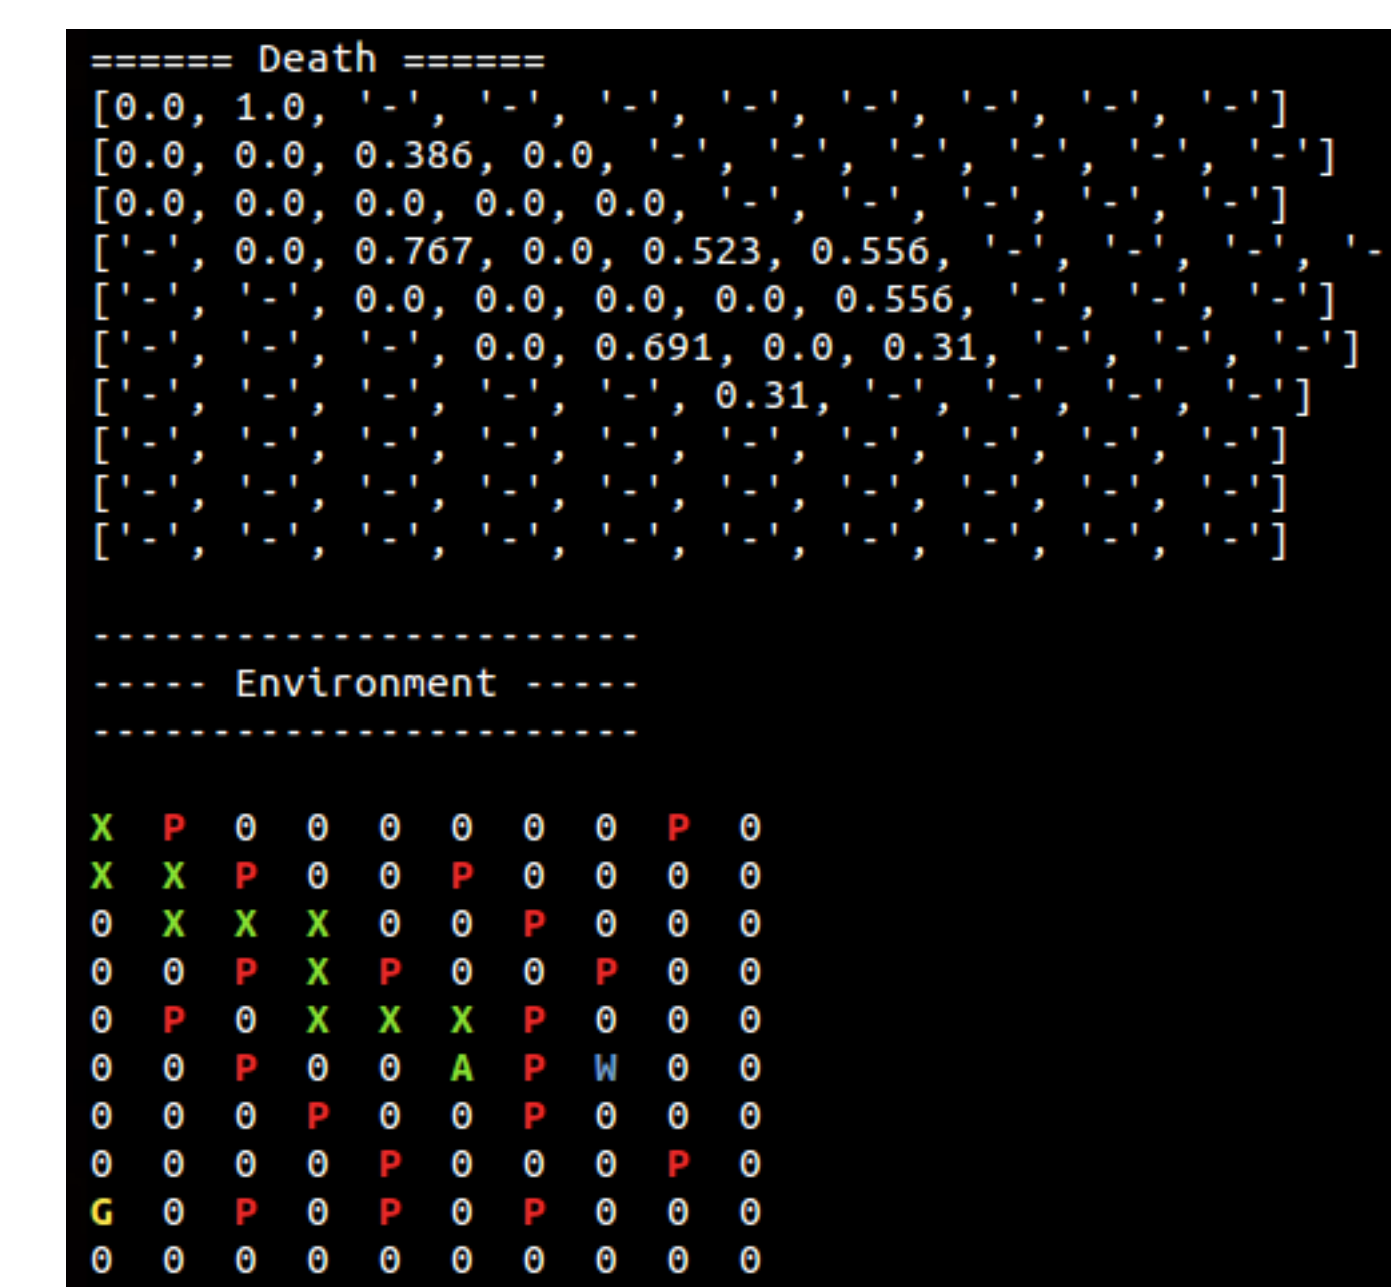


Figure 4 : A 10×10 WW environment

Pit Probability

Calculating $P(p_{\mathbf{x}})$, the probability that a pit is located at square \mathbf{x} , is a much more difficult task.

During gameplay, the agent builds a set of models (one for each possible world), and then selectively filters this set as the game goes on. Essentially, it can be shown that $P(p_{\mathbf{x}}) = \sum_i^n P(M_{i\mathbf{x}})$ for all n models where \mathbf{x} contains a pit.

Important Note

To optimize this process, the agent only models worlds for the squares contained in V' . This set includes all possible "next move" squares. Otherwise, a 10×10 environment like that shown in Figure 4 would require 2^{100} models be built!

Conclusion

A 5000 trial test was run on this implementation of the wumpus world agent (using a 4×4 environment). The results were as follows:

- Found Gold: 2568 (51%)
- Flopped: 2013 (40%)
- Death: 237 (4%)
- Forfeited: 182 (3%)

I have defined "flop" in this context to indicate that the agent forfeited on the first move (as it will if it spawns directly adjacent to a pit or wumpus).

Therefore, of the environments that did not result in a flop, the agent found the gold 86% of the time, died 8% of the time, and forfeited 6% of the time.

References

- [1] Stuart Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Prentice Hall, 2009
- [2] Sheldon Ross. *A First Course in Probability*. Prentice Hall, 2009
- [3] Kenneth H. Rosen. *Discrete Math and Its Applications*. McGraw-Hill, 2012

Contact Information

- Web: bryanbugyi.com
- Email: bryan_bugyi@mymail.rcbc.edu

The source code for this project can be found on my GitHub page at github.com/bbugyi200/WumpusWorld.